


Ray Tracing

CSU44052 Computer Graphics




Binh-Son Hua

Rendering

- Rendering is fundamentally concerned with determining the most appropriate colour (i.e. RGB tuple) to assign to a pixel associated with an object in a scene.




Global Illumination



Ray Tracing


Sharp
shadows

Perfectly
specular
reflections



Phong
Illumination

Ray Tracing is a View-Dependent Solution




Scene Geometry



Solution determined only for directions through pixels in the viewport



First ray-traced image




Turner Whitted, 1979

Ray-tracing history

- First used in computer graphics in 1980
- Integrated reflection, refraction, hidden surface removal, shadows in a single model
- Rays usually considered to be infinitely thin
 - Reflection & refraction occur without any spreading
 - Perfectly smooth surfaces
 - Not real-world – like a wall of mirrors
- “Super-real” images at a high cost



Forward Ray tracing




- Only a fraction of rays reach the image
- Many, many rays are required to get a value for each pixel.


Backward Ray Tracing

For each pixel in the viewport:

- Trace a ray from the eye (called the eye ray) into the scene, through the pixel.
- Determine the first object hit by the ray \Rightarrow ray casting.
- We then shade this using an extended form of the Phong model.



Ray Tracing



Ray Tracing

- The eye ray will typically intersect a number of objects, some more than once
⇒ sort intersections to find the closest one.
- The Phong illumination model is then evaluated BUT:
 - we trace a reflected ray if the surface is specular
 - we trace a refracted ray if the surface is transparent
 - we trace shadow rays towards the light sources to determine which sources are visible to the point being shaded
- The reflected/refracted rays themselves will hit surfaces and we will **recursively** evaluate the illumination at these points.
- A very large number of rays must be traced to illuminate a single pixel.

Whitted Illumination Model


- Whitted Illumination Model

$$L_r(x, V) = \underbrace{L_{emitted}(x, V) + L_{Phong}(x, V)}_{\text{local contribution}} + \underbrace{L_{reflected}(x, V) + L_{refracted}(x, V)}_{\text{global contribution}}$$


- Ray tracing is a hybrid local/global illumination algorithm
- We only consider global lighting effects from ideal specular directions
- Also, before adding each light's Phong contribution we determine if it is visible to point x, thus allowing shadows to be determined.

Recursive Ray Tracing

- The ray tracing algorithm is recursive, just as the radiance equation is recursive.
- At each intersection we trace a specularly reflected and transmitted ray (if the surface is specular) or terminate the ray if diffuse.
- Thus we trace a ray back in time to determine its history, beginning with the eye ray: this leads to a binary tree




Ray Tree



Recursive Ray Tracing

- In theory, this recursive process could continue indefinitely.
- In practice, at each intersection the ray loses some of its contribution to the pixel (i.e. its importance decreases).
 - if the eye ray hits a specularly reflecting surface with reflectivity of 50%, then only 50% of the energy hitting the surface from the reflected direction is reflected towards the pixel.
 - if the next surface hit is of the same material, the reflected ray will have its contribution reduced to 25%.
- We terminate the recursion if:
 - the current recursive depth $>$ a pre-determined maximum depth or
 - if the ray's contribution to the pixel $<$ some pre-determined threshold ϵ

Recursion Clipping



Very high maximum recursion level



Max level = 1



Max level = 2



Max level = 3




Max level = 4

The Ray

- Mathematically, a ray is the *affine half-space* defined by:

$$\mathbf{r} = O + t\vec{d} \quad t \geq 0$$




- All points on the ray correspond to some positive value of t , the *parametric distance* along the ray. If \mathbf{d} is *normalised* then t is the length along the ray of the point.

Ray Tracing Algorithm

```
for each pixel in viewport
{
    determine eye ray for pixel
    intersection = trace(ray, objects)
    colour = shade(ray, intersection)
}
```

```
trace(ray, objects)
{
    for each object in scene
        intersect(ray, object)
    sort intersections
    return closest intersection
}
```



Ray Tracing Algorithm

```
colour shade(ray, intersection)
{
    if no intersection
        return background colour


    for each light source
        if(visible)
            colour += Phong contribution

    if(recursion level < maxlevel and surface not diffuse)
    {
        ray = reflected ray
        intersection = trace(ray, objects)
        colour +=  $\rho_{\text{refl}}$  * shade(ray, intersection)
    }
    return colour
}
```

Ray Casting


```
for each pixel in viewport
{
    determine eye ray for pixel
    intersection = trace(ray, objects)
    colour = shade(ray, intersection)
}
```

```
trace(ray, objects)
{
    for each object in scene
        intersect(ray, object)
    sort intersections
    return closest intersection
}
```




Ray Casting

- For each sample
 - Construct ray from eye position through view plane
 - Find first surface intersected by ray through pixel
 - Compute colour sample based on surface radiance



Ray Casting


- Collision detection
- Object picking
- Line of sight
- Occlusion culling



The Ray

- Mathematically, a ray is the *affine half-space* defined by:


$$\mathbf{r} = O + t\vec{d} \quad t \geq 0$$



- All points on the ray correspond to some positive value of t , the *parametric distance* along the ray. If \mathbf{d} is *normalised* then t is the length along the ray of the point.

Ray-Object Intersection Testing

- Once we've constructed the eye rays we need to determine the intersections of these rays and the objects in the scene.
- Upon intersection we need the **normal** vector at the point of intersection in order to perform shading calculations.




The Sphere

- A sphere of center (C_x, C_y, C_z) with radius r is given by:

$$f(x, y, z) = (x - C_x)^2 + (y - C_y)^2 + (z - C_z)^2 - r^2 = 0$$

- Question: is the point $(5, 1, 0)$ on this sphere?



The Sphere

- A sphere object is defined by its center C and its radius r .

- Implicit Form:
$$f(\vec{v}) = |\vec{v} - C|^2 - r^2 = 0$$
$$f(x, y, z) = (x - C_x)^2 + (y - C_y)^2 + (z - C_z)^2 - r^2 = 0$$

- Explicit Form:
$$x = f_x(\theta, \phi) = C_x + r \sin \theta \cos \phi$$
$$y = f_y(\theta, \phi) = C_y + r \cos \theta$$
$$z = f_z(\theta, \phi) = C_z + r \sin \theta \sin \phi$$

- We can use either form to determine the intersection; we will choose the implicit form.

Ray Sphere Intersection

- All points on the ray are of the form: $\mathbf{ray} = O + t\vec{d} \quad t \geq 0$
- All points on the sphere satisfy:

$$(x - C_x)^2 + (y - C_y)^2 + (z - C_z)^2 - r^2 = 0$$

- Any intersection points (= points shared by both) must satisfy both, so substitute the ray equation into the sphere equation and solve for t:

$$\underbrace{([O_x + td_x] - C_x)^2}_{\substack{\uparrow \\ \text{ray equation}}} + \underbrace{([O_y + td_y] - C_y)^2}_{\substack{\uparrow \\ \text{ray equation}}} + \underbrace{([O_z + td_z] - C_z)^2}_{\substack{\uparrow \\ \text{ray equation}}} - r^2 = 0$$

Problem

$$\underbrace{([o_x + td_x] - c_x)^2}_{\uparrow} + \underbrace{([o_y + td_y] - c_y)^2}_{\uparrow} + \underbrace{([o_z + td_z] - c_z)^2}_{\uparrow} - r^2 = 0$$

ray equation

- Expand the first term
 - remember $(a-b)^2 = a^2 - 2ab + b^2$

- Rearrange into: $At^2 + Bt + C = 0$

Ray Sphere Intersection

- Rearrange and solving for t leads to a quadratic form (which is to be expected as the sphere is a quadratic surface):

$$At^2 + Bt + C = 0$$

$$A = (d_x^2 + d_y^2 + d_z^2) = 1$$

$$B = 2d_x(O_x - C_x) + 2d_y(O_y - C_y) + 2d_z(O_z - C_z)$$

$$C = (O_x - C_x)^2 + (O_y - C_y)^2 + (O_z - C_z)^2 - r^2$$

- We employ the classic quadratic formula to determine the 2 possible values of t:

$$t = \frac{-B \pm \sqrt{B^2 - 4AC}}{2A} = \frac{-B \pm \sqrt{B^2 - 4C}}{2}$$

Intersection Classification

- Depending on the number of real roots we have a number of outcomes which have nice geometric interpretations
- We use the discriminant: $d = B^2 - 4C$


$d = 0 \Rightarrow$ 1 root (ray is tangent to sphere)

$d < 0 \Rightarrow$ no real roots (ray misses)

$d > 0 \Rightarrow$ 2 real roots:

Both positive \Rightarrow sphere in front of ray

One negative \Rightarrow ray origin inside sphere



Ray Sphere Intersection Test

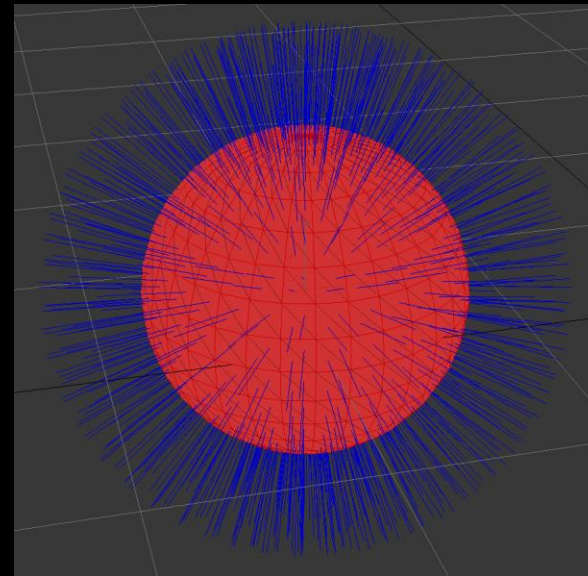

- If we have 2 positive values of t we use the **smallest** (i.e. the nearest to the origin of the ray).
- t is substituted back into the ray equation yielding the point of intersection: $P_{int} = O_{ray} + t_{int}d_{ray}$
- We then evaluate the Phong model at this point. To do so we need the normal to the surface of the sphere at the point of intersection.
- The normal and the original ray direction are then used to determine the directions of reflected and refracted rays.

Normal to Sphere


- We can compute the normal to a sphere with centre C at a point x as:

$$N = \frac{\vec{x} - C}{|\vec{x} - C|}$$

i.e. the normal to the sphere is the normalised vector associated with the point.



Ray Polygon Intersection




To test if the ray intersects the polygon:

- Assume planar polygons
- First see if the ray intersects the polygon's plane
- If it does, see if the intersection point is inside or outside the polygon

Plane Definition

- A plane may be defined by its normal and the perpendicular distance of the plane to the origin:



Equation of a Plane

$$\vec{v} \cdot \vec{n}_P = d_p \quad \text{or}$$

$$n_x x + n_y y + n_z z - d_P = 0$$

Ray Plane Intersection

1. Does ray intersect the polygon's plane?

- First compute the dot product between the normalised ray and the polygon's normal (i.e., the plane's normal):
 - If dot product == 0 it => ray & normal perpendicular
 - No intersection!
 - If > 0 => ray and normal are in the same direction
 - Back face intersection
 - if < 0 => plane facing direction of ray
 - There is an intersection!
- Note there can only be one intersection. The normal to the plane at each point is the same, i.e., the plane normal \vec{n}_p

$$\vec{n}_p \cdot \vec{d}_r$$

Ray Plane Intersection

2. Now find intersection point of ray with the plane

To intersect a ray with a plane we substitute in the ray equation and solve for t

$$(O_r + td_r) \cdot \vec{n}_p = d_p$$

- O is origin of the ray and d_r is the direction of the ray. Solve for t and that will give you the point on the plane

$$\Rightarrow t = \frac{d_p - \vec{n}_p \cdot O_r}{\vec{n}_p \cdot \vec{d}_r}$$

- t is substituted back into the ray equation yielding the point of intersection:


$$P_{int} = O_{ray} + t_{int}d_{ray}$$

Ray Polygon Intersections

3. Next, determine if the intersection point is within the polygon interior.

The *Jordan Curve Theorem*:

- construct any ray with the intersection point as an origin
- count the number of polygon edges this ray crosses (= *winding number*)
- if w is odd then the point is in the interior



Ray Triangle Intersections

3. Next, determine if the intersection point is within the **triangle** interior.


Use barycentric coordinates!

$$P = aA + bB + cC$$

with the constraint

$$a + b + c = 1$$

If $0 \leq a, b, c \leq 1$: point P is in the triangle.




Secondary Rays

- Having determined the closest intersection point x along a ray path we then evaluate the *Whitted illumination model* at this point:
 - Construct a shadow ray to each light source
 - Compute the local illumination at x
 - Construct a *reflected* ray from x and recurse
 - Construct a *refracted* ray from x and recurse


Secondary Rays

N , I , R and T all lie in the same plane




Reflection

- The law of reflection says that for specular reflection, the angle at which the wave is incident on the surface equals the angle at which it is reflected.



Reflected Ray



$$\begin{aligned} R &= I + 2N(-I \cdot N) \\ &= I - 2N(I \cdot N) \end{aligned}$$

$(-I \cdot N)$ is the length of I projected onto N

Refraction


- Refraction is the bending of the path of a light wave as it passes from one material to another material.
- The refraction occurs at the boundary and is caused by a change in the speed of the light wave upon crossing the boundary.
- The tendency of a ray of light to bend one direction or another is dependent upon whether the light wave speeds up or slows down upon crossing the boundary.



Credit: Google


Refraction

- A comparison of the angle of refraction to the angle of incidence provides a good measure of the refractive ability of any given boundary.



Refracted Rays

- We use a geometric construction using *Snell's Law* to determine \mathbf{T} :




$$\eta = \frac{\eta_i}{\eta_t} = \frac{\sin \theta_t}{\sin \theta_i} = \frac{|\vec{v}|}{|\vec{u}|}$$

$$\mathbf{T} = \vec{w} + \vec{v} = \vec{w} + \frac{|\vec{v}|}{|\vec{u}|} \vec{u} = \vec{w} + \eta \vec{u}$$

Total Internal Reflection

- *Total internal reflection* occurs when the incident angle exceeds the *critical angle* for the surface. This will only happen when passing from a *higher refractive index* to a *lower one*, e.g., from water to air.




Critical angle can be estimated when angle of refraction is 90 degrees, so $\frac{\eta_i}{\eta_t} = \frac{1}{\sin \theta_i}$


Shadows

- Whether a point is in shadow or not is determined by casting a ray from each intersection point to the light source
- If it intersects any object then the point of interest is deemed to be in shadow
 - Easier than ray/object intersections, as we only need to know if intersection has occurred (do not need to find the nearest object)
- Shadow calculations impose a computational overhead in ray tracing that increases rapidly as the number of light sources increases.

Soft Shadows




Point Light Source



- A “point” light source is just a convenient model, and only allows for hard/sharp shadows.
- Consider a light bulb, for example. It is not an infinitesimally small point. It has volume.
- Implication: Real light sources produce soft shadows.


Soft Shadows

- Shadows are not uniformly dark. The shadow is divided into two parts: the **umbra** and **penumbra**:
- **Umbra**: no light at all from the light source. Completely dark.
- **Penumbra**: some light from the light source. Partially dark.




Soft Shadows in Ray Tracing

- For more realistic shadows, model light sources as **volumes** rather than points.
- Use a sphere to model a light source, rather than a point.
- This is often quite realistic because many light sources in real life are spherical, e.g. light bulbs and lanterns.
- When calculating lighting, shoot **several** rays to the light source, instead of just one ray.
- Calculate lighting for each ray, and take the average.



Point light source: The surface is completely lit by the light source.



Finite light source: 3/5 of the rays reach the light source. The surface is partially lighted.

Further Reading

- Physically Based Rendering: <https://pbr-book.org/>

Chapter 1.2: Photorealistic Rendering and the Ray-Tracing Algorithm.

Chapter 13: Light Transport I: Surface Reflection.